

8: Estimation II

DEPARTMENT OF MECHANICAL ENGINEERING

Mobile Robotics (ENME 650)

Instructor:


DR. ALEX RAMIREZ-SERRANO

Office: MEB 523

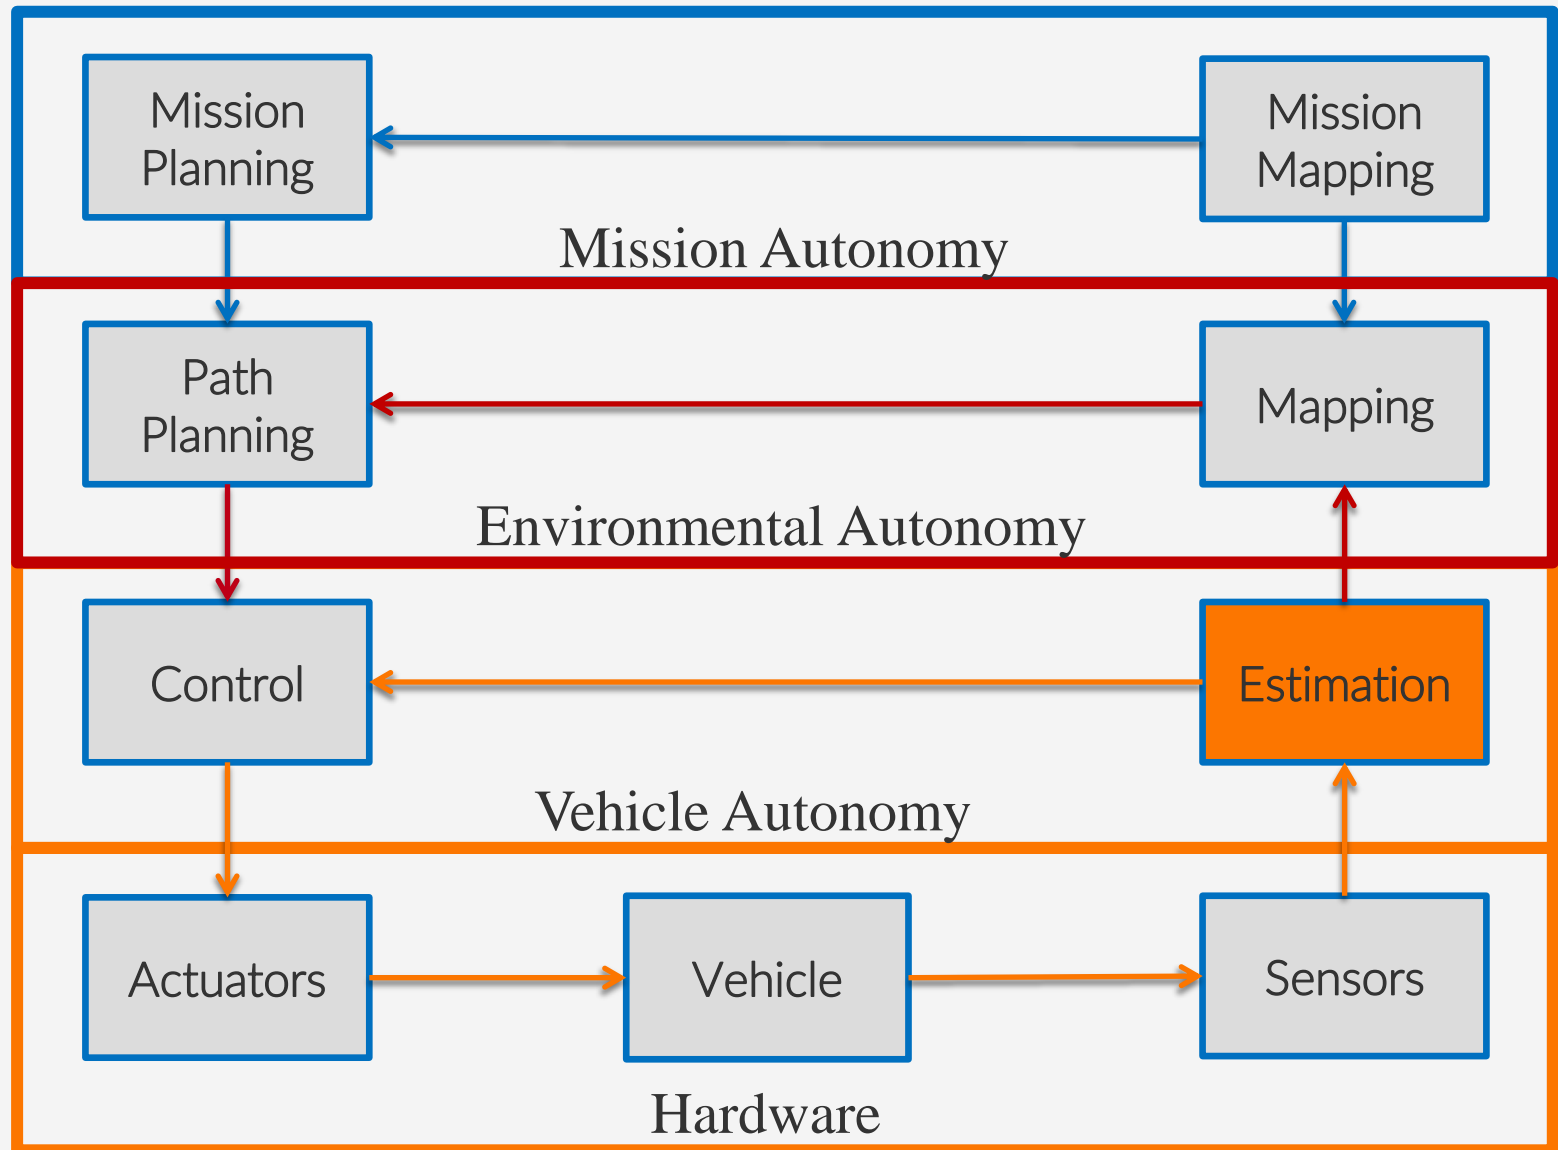
Phone: 220-3632

E-mail: aramirez@ucalgary.ca

CANADA



COMPONENTS



OUTLINE

- Review of Probability
- Bayes Filter Framework
- Kalman Filter
- Extended Kalman Filter
- Unscented Particle Filter
- Particle Filter

PARTICLE FILTERS

- The Bayes Filter Framework has now been adapted to:
 - Kalman Filter (KF)
 - *Linear models with additive Gaussian Noise*
 - Extended Kalman Filter (EKF)
 - *Nonlinear models with additive Gaussian noise*
- Both continuous Gaussian methods are computationally appealing (*i.e., can run in real-time*)
 - Even for large numbers of state & measurement variables
 - Benefit arises from ability to maintain Gaussian beliefs
 - *Track only mean and covariance throughout filtering process*

These filters are called **Parametric filters** because we have parametrized the distribution with a mean and a covariance.

PARTICLE FILTERS

- In both cases, modeling requirements rule out a significant portion of real systems
 - *Nonlinear systems where linearization is a poor approximation over distribution range*
 - *Systems with multiple reasonable hypotheses*
- Alternatives include **non-parametric filters** (e.g., when we cannot assume gaussian distributions, doesn't make sense to linearize models, etc.):
 - Filters that do not track distribution parameters
 - Bayes/Histogram Filter
 - *Discrete state systems with known probabilities*
 - *Explodes computationally for higher dimensional models*
 - **Particle Filter**
 - *Maintain a sample set representation of beliefs*
 - *Results can be poor in higher dimensional models*
 - *Also called Sequential Monte Carlo methods*

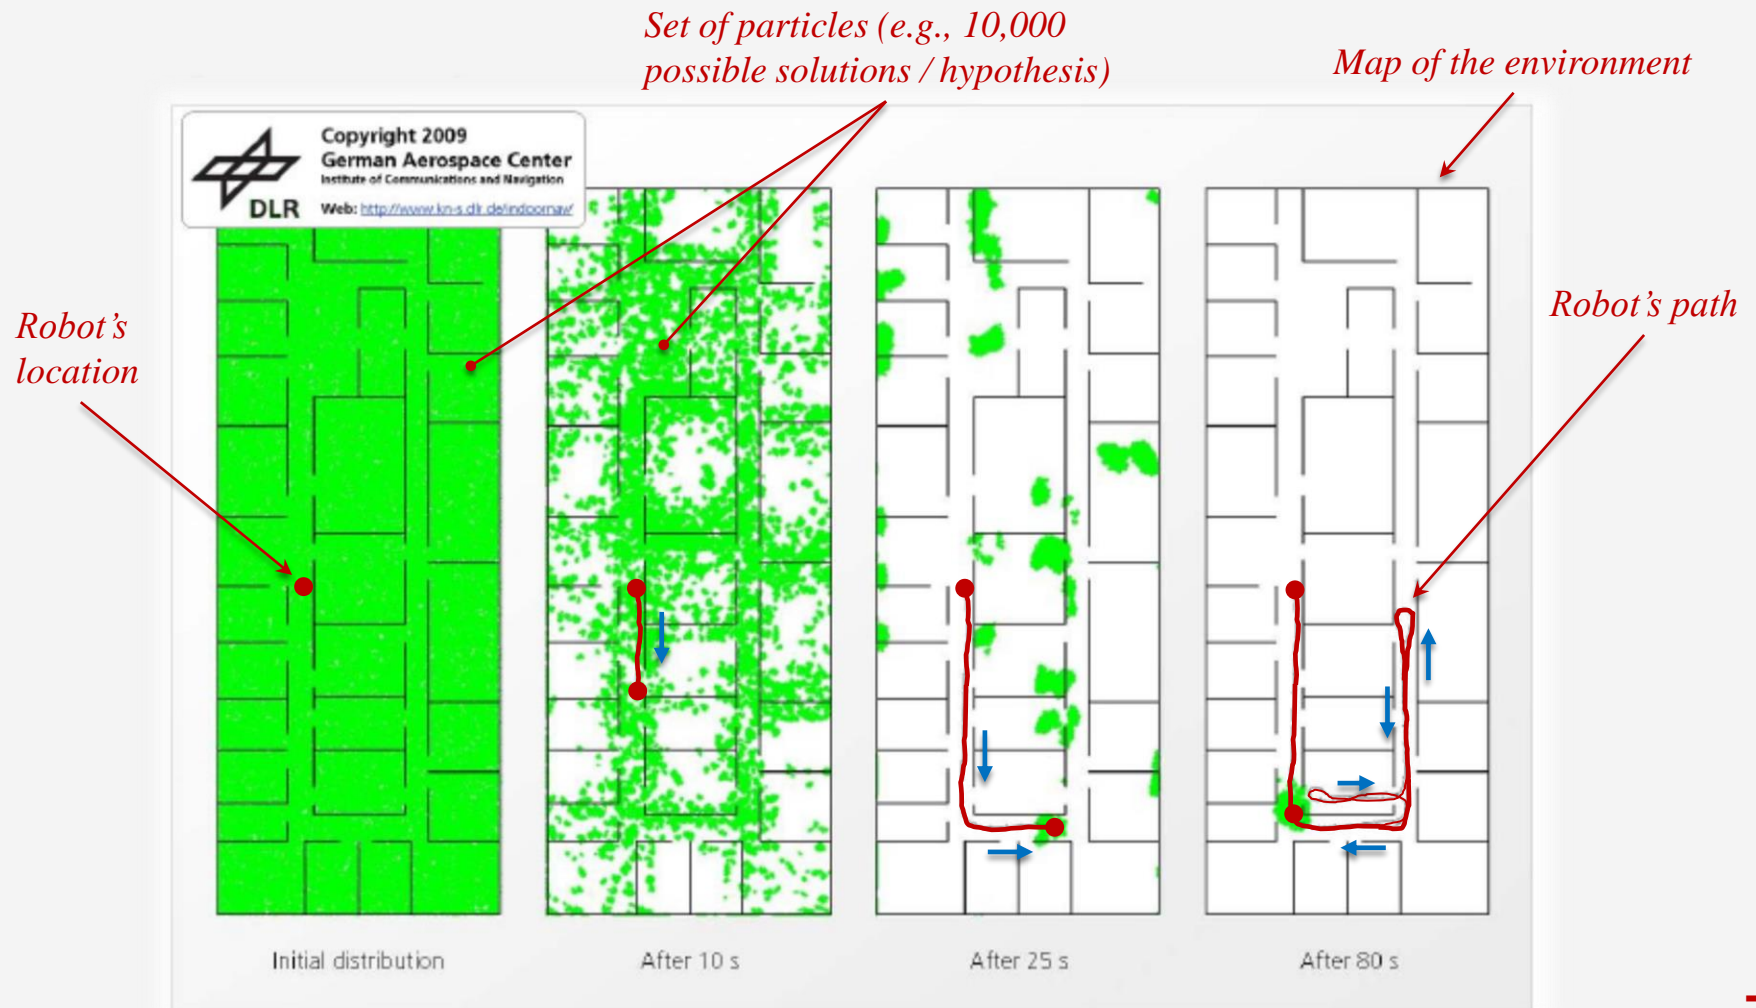
PARTICLE FILTERS

- Modeling assumption(s):

- Instead of assuming Gaussian, tracking μ_t, Σ_t , generate a set of sample states from each distribution
 - Each sample is a **hypothesis** about the current state
 - Properties of the whole collection of samples are used to generate estimates
- Not possible to sample belief distribution directly, **must apply Importance Sampling** (*to find a belief from the sample set that indicates what the solution is - where the robot is*)

PARTICLE FILTERS

- Example particle sets – density of points defines probability



PARTICLE FILTERS

- Generating samples from a known distribution
 - Given a probability density function, draw samples with the appropriate probability
 - Easy for uniform, Gaussian
 - Use built in Matlab functions
 - Harder for arbitrary pdfs ([probability distribution functions](#)), but approximation is possible
 - Needed in Particle filters to perform measurement update


PARTICLE FILTERS

- Generating samples from a distribution (*Steps to follow*):
 - Given a state $x \in \mathbb{R}$ and a distribution $g(x): \mathbb{R} \rightarrow [0,1]$
 1. Create a vector \mathbf{X} of evenly spaced values of x over the range of interest
 - e.g., If $g(x)$ is Gaussian create \mathbf{X} to span $\pm 5\sigma$ about mean
 2. Create an exact/approximate cumulative distribution vector, $G(\mathbf{X})$
 - Integrate probability distribution $g(x)$ to get a cumulative $G(x)$, and create the vector $G(\mathbf{X})$
 - Or sum probabilities $g(\mathbf{X})$ and normalize to get vector $G(\mathbf{X})$
 3. Draw samples from a uniform distribution over $[0,1]$
 4. Find closest value to sample in $G(\mathbf{X})$
 5. Corresponding value of x is a sample, denoted $x_g^{[i]}$

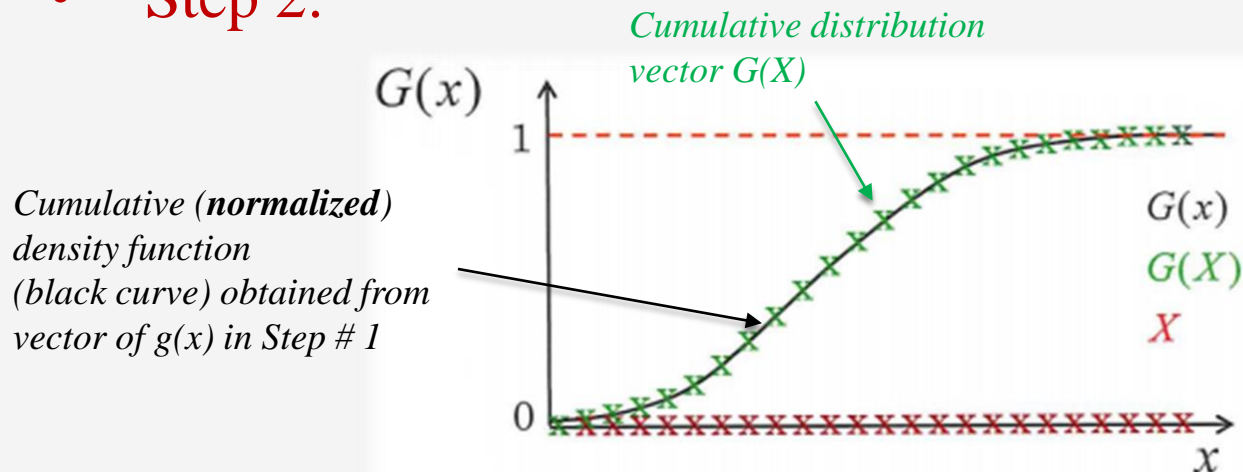
PARTICLE FILTERS

- Sampling of $g(x)$ (*graphical representation of steps in previous slide*)

- **Step 1:**



- **Step 2:**



PARTICLE FILTERS


○ Sampling of $g(x)$

• Steps 3 to 5:

Values having **HIGH**
probability of being
selected

Cumulative (*normalized*)
density function

STEP 3:
Randomly
selected value
between 0 and 1



Values having **LOW**
probability of being
selected

STEP 5:

Select the closest sample in the
vector **X** that corresponds to the
randomly selected value, $u^{[i]}$

Good sample (particle) that
describes the hypothesis

PARTICLE FILTERS

○ Importance Sampling

- However, we do not know $f(x)$, so assuming we have another distribution, $g(x)$, that is directly related to $f(x)$, we can use such function to find $f(x)$.
- **Goal:** perform a calculation using a distribution, $f(x)$, but without being able to sample it directly
 - $f(x)$ = Target distribution, unknown
- Can first sample a different distribution, $g(x)$,
 - $g(x)$ = Proposal distribution, known

$$f(x) = F(g(x))$$


$$f(x) = \frac{f(x)}{g(x)} g(x)$$

Sampling
function

$g(x)$

Resampling function
(incorporating new
information)

$f(x)$



PARTICLE FILTERS

- Importance Sampling

- Then use relationship between distribution if known to define the weighting factor as

$$w(x) = \frac{f(x)}{g(x)}$$

- Finally, resample from $g(x)$, with weights $w(x)$ to generate samples of $f(x)$
 - If weighting factor is known, can perform this calculation without knowing $f(x)$
 - Note that $g(x) > 0$ wherever $f(x) > 0$ for this to be valid

PARTICLE FILTERS

- Importance Sampling

- Define weights for each sample $x_g^{[i]}$ in S


$$w(x) = \frac{f(x)}{g(x)}$$

- Weights are the probability that we should include sample $x_g^{[i]}$ in our final sample set
 - The importance of sample $x_g^{[i]}$
- Not obvious how to calculate the weight at this point, will become clear in derivation of particle filter
- For now, found by dividing $f(x_g^{[i]})$ by $g(x_g^{[i]})$
 - This assumes complete knowledge of $f(x)$,
 - ❖ (yes, cheating)

PARTICLE FILTERS

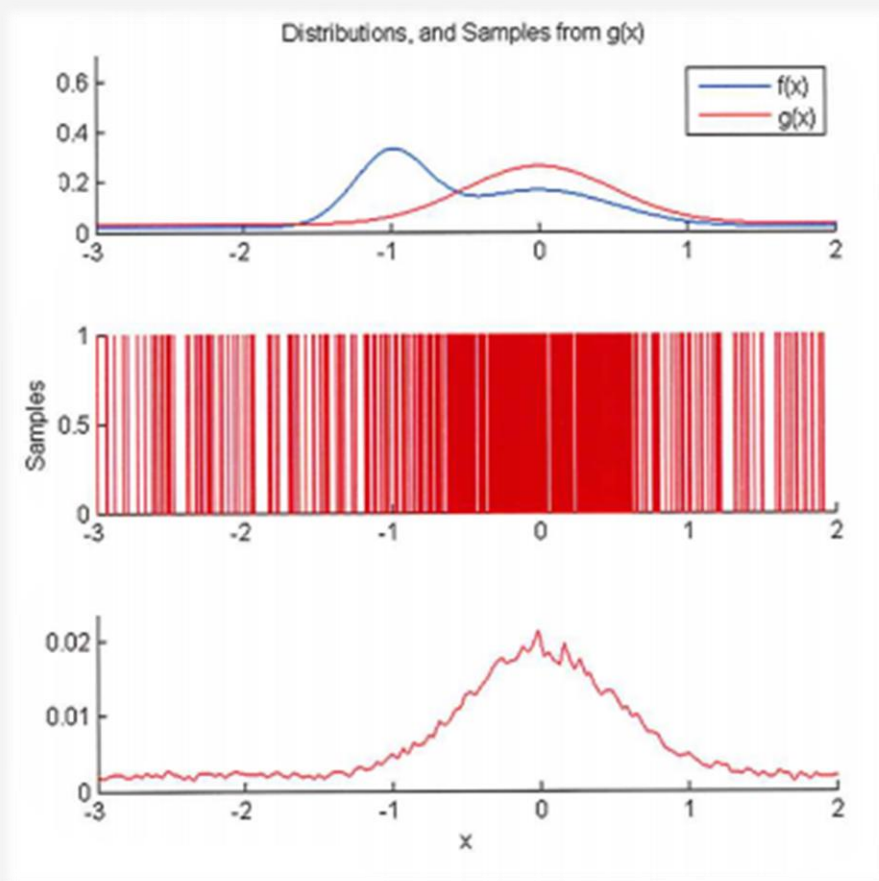
○ Importance Sampling

- Importance sampling of $f(x)$
 1. Define cumulative distribution $W(x)$ based on weights $w(x)$ as before (*samples need not be ordered*)
 2. For each sample:
 - a) Take uniform sample, $u^{[i]}$
 - b) Find first element of $W(x)$ that exceeds current sample
 - c) Add corresponding value of $x_g^{[i]}$ as a sample-to-sample set, S'



PARTICLE FILTERS

- Example (*Theoretical example: we know $f(x)$ and $g(x)$*)
 - Target distribution $f(x)$, proposal distribution $g(x)$
 - 20000 samples drawn from $g(x)$ ($1/25^{th}$ of samples shown)

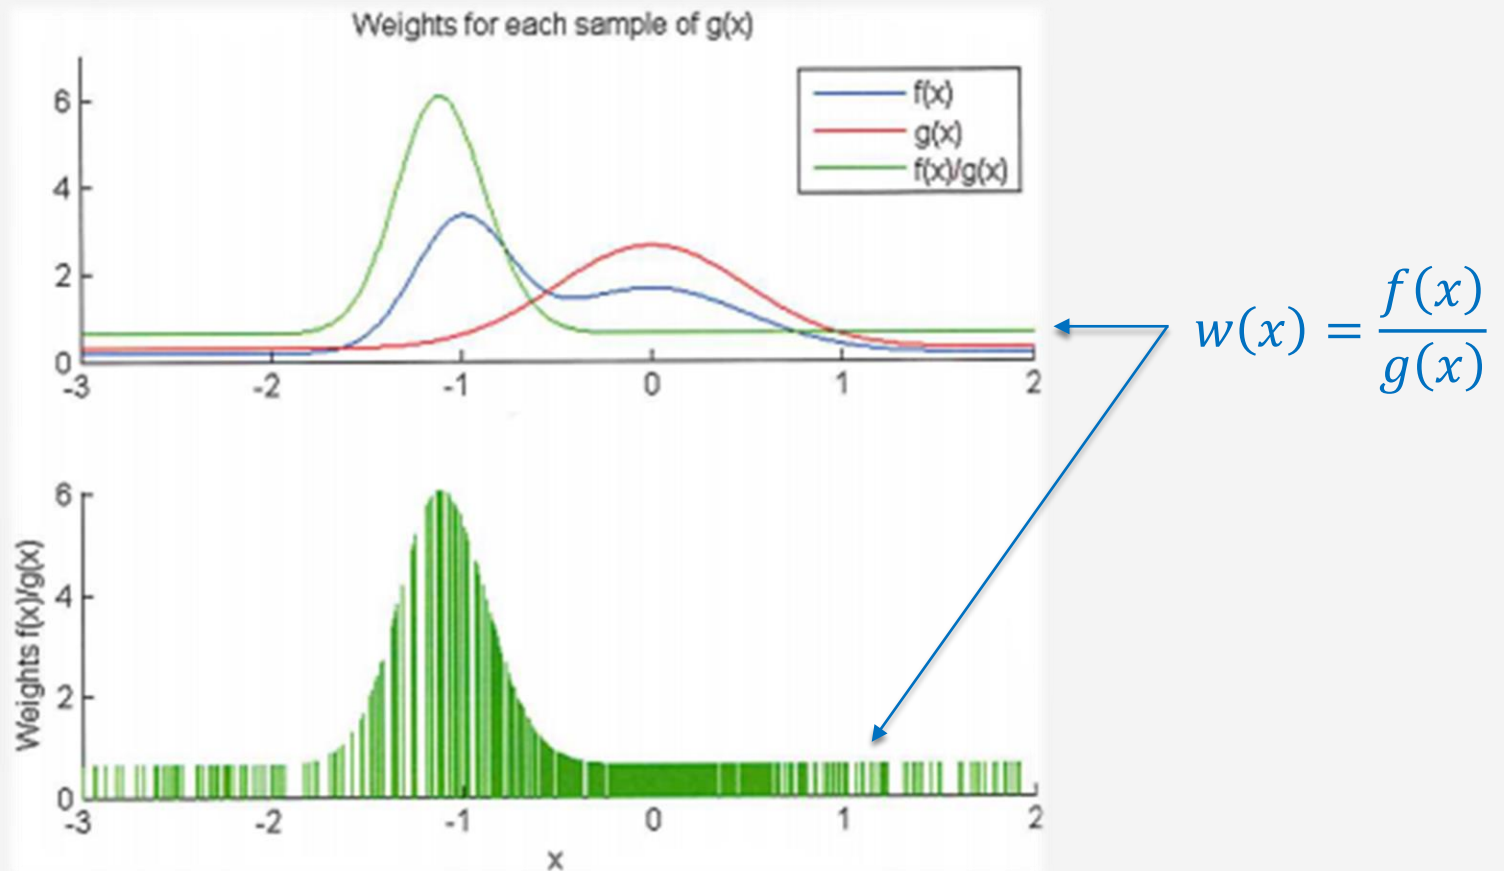


← We satisfy $g(x) > 0$
wherever $f(x) > 0$

We want to reconstruct
 $f(x)$ from $g(x)$

PARTICLE FILTERS

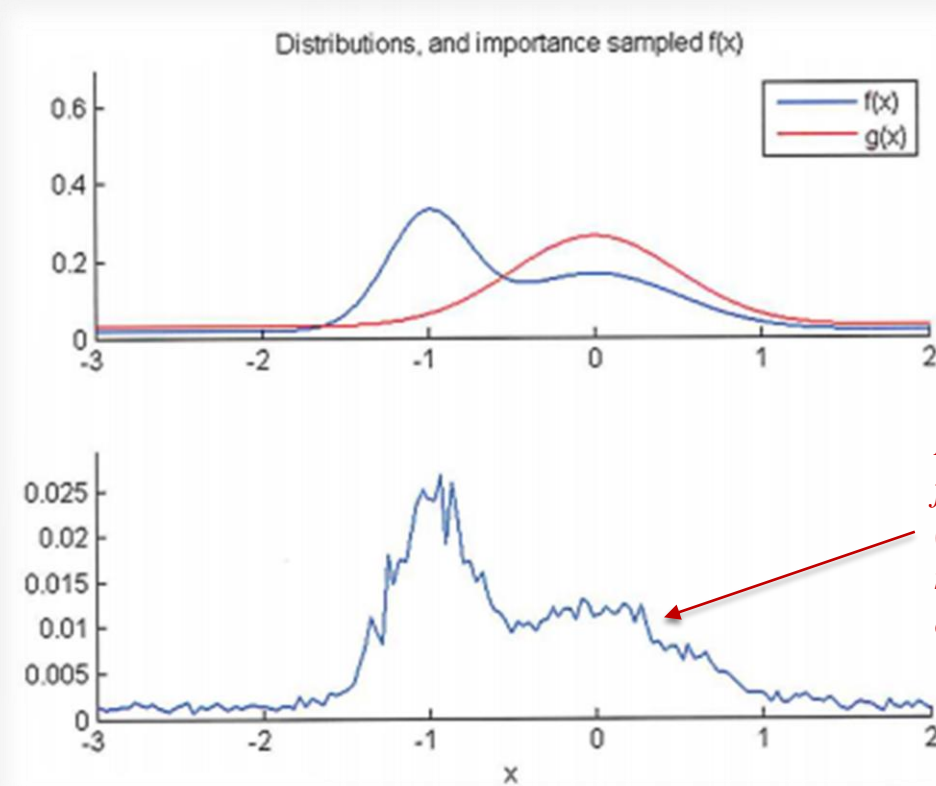
- Example
 - We now calculate the weights for all x and sample weights for each sample



PARTICLE FILTERS

○ Example

- Importance sample points to generate new set
- New set is distributed according to $f(x)$



New distribution obtained by forming the cumulative (normalized) density function and making our selection from there to construct $f(x)$

PARTICLE FILTERS

○ The Particle Set

- A sample can be drawn from a *proposal distribution*

$$x_p^{[i]}$$

- The sample is assigned a weight

$$w_p^{[i]}$$

- The combination of sample and weight is a particle

$$s^{[i]} = \{x_p^{[i]}, w_p^{[i]}\}$$

- The particle set is used to generate an approximation to the *target distribution*

$$s^{[i]} = \{x^{[i]}, w^{[i]}\} \quad S = \{s^{[1]}, \dots, s^{[I]}\}$$

- I is the total number of particles in the set
- The approximation improves as $I \rightarrow \infty$

PARTICLE FILTERS

- Defining the usual model elements, in general probabilistic form (*Derivation*)

- State prior:

$$p(x_0)$$

- Motion Model:

$$p(x_t | x_{t-1}, u_t)$$

- Measurement Model:

$$p(y_t | x_t)$$

- Only restriction on model elements are that samples can be drawn from them, (*probabilities known for all conditional values*)

PARTICLE FILTERS

○ Beliefs

- In particle filters, the belief distributions will be represented by particle sets

- The belief

$$x_t^{[i]} \sim \text{bel}(x_t) = p(x_t | y_{1:t}, u_{1:t})$$

$$S_t = \{s_t^{[1]}, \dots, s_t^{[I]}\}$$

- The predicted belief

$$\bar{x}_t^{[i]} \sim \overline{\text{bel}}(x_t) = p(x_t | y_{1:t-1}, u_{1:t})$$

$$\bar{S}_t = \{\bar{s}_t^{[1]}, \dots, \bar{s}_t^{[I]}\}$$

PARTICLE FILTERS

- Particle Filter Algorithm:

- 1. Prediction Transformation

- Transform prior belief particle set to predicted belief through sampling

- 2. Importance factor

- Using measurement, calculate particle importance factor
 - *Probability of the measurement occurring, given the state was defined by the current particle*

- 3. Resampling

- Transform predicted belief particle set to belief using importance sampling

Note: Steps 1, 2 can be combined into a single loop, if prediction and measurement steps are combined.

PARTICLE FILTERS

○ Particle Filter Components

1. Prediction Update

- The samples $x_{t-1}^{[i]}$ are known from previous iteration
- The motion model and input are known
- It is therefore possible to generate samples of $\overline{bel}(x_t)$

$$\bar{x}_t^{[i]} \sim p(x_t | x_{t-1}^{[i]}, u_t)$$

- One new sample is drawn from each distribution defined by the prior samples
- The set of I new samples defines an approximation to $\overline{bel}(x_t)$
- Unit weighting on each particle

$$\bar{S}_t = \{\bar{s}_t^{[1]}, \dots, \bar{s}_t^{[I]}\}$$

PARTICLE FILTERS

○ Particle Filter Components

2. Measurement update

- The measurement is known but the state is not
- Would like to generate a particle set to represent $bel(x_t)$
 - Target distribution
- Have particle set representation of prediction belief $\overline{bel}(x_t)$
 - Proposal distribution
- Use importance sampling to generate belief update

measurement model \rightarrow $w_i^{[i]} = \eta \frac{bel(x_t)}{\overline{bel}(x_t)}$ \leftarrow prior belief

- The proper weighting to use turns out to be

$$w_t^{[i]} = p(y_t | \bar{x}_t^{[i]})$$

PARTICLE FILTERS

○ Particle Filter Expanded Algorithm

1. Prediction update:

a) Sample $\bar{x}_t^{[i]} \sim p(x_t | x_{t-1}^{[i]}, u_t)$

b) Weight $\bar{w}_t^{[i]} = 1$

c) Add to \bar{S}_t

Prior belief

2. Measurement update:

1. For each particle in \bar{S}_t

a) Calculate weighting


$$w_t^{[i]} = p(y_t | \bar{x}_t^{[i]})$$

2. For $j = 1$ to I

a) Draw particle $\bar{x}_t^{[i]}$ with probability $\propto w_t^{[i]}$

b) Add to S_t as $s_t^{[i]} = \{x_t^{[i]}, 1\}$

Belief



PARTICLE FILTERS

- Particle Filter Expanded Algorithm (*simplified*)

1. For each particle in S_{t-1}

- a) Propagate sample forward using motion model (*sampling*)

$$\bar{x}_t^{[i]} \sim p(x_t | x_{t-1}^{[i]}, u_t)$$

- b) Calculate weight (*importance*)

$$w_t^{[i]} = p(y_t | \bar{x}_t^{[i]})$$

- b) Store in interim particle set

$$S'_t = S'_t + \{s_t^{[i]}\}$$

2. Normalize weights

3. For $j = 1$ to I

- a) Draw index i with probability $\propto w_t^{[i]}$

- Add to final particle set

$$S_t = S_t + \{s_t^{[i]}\}$$

(*resampling*)



PARTICLE FILTERS


○ Example:

- Robot Localization

- Robot travels along a hallway, can detect doors within a range with a noisy sensor
- Knows probability of detecting a door, given a specific location


$$p(y_t | x_t)$$

- Knows motion model, and has uniform initial belief (*has no idea where it is*)



PARTICLE FILTERS

- Example
 - **Step One:**
 - Sample uniform over state space




PARTICLE FILTERS

- Example

- Step Two:

- Propagate samples through motion model (*i.e., provide an input, u_t , and add disturbances*)




PARTICLE FILTERS

- Example

- Step Three:

- Take a measurement, and use $p(y_t|x_t)$ to calculate weights




- Particles that are more likely have higher weights
 - Starting to narrow down options
 - Still difficult to estimate state (mean?)

PARTICLE FILTERS

- Example


- **Step Four:**
 - Perform **resampling** to get more particles in areas of higher probability
 - Reset weights to 1, as particle locations capture probability information
- Repeat
- The following particle set shows how the motion model distributes the identical particles that result from resampling



PARTICLE FILTERS

○ Example

- After a second measurement, weights are again assigned to the particles




- True state starts to become apparent

PARTICLE FILTERS

○ Example

- After resampling again, propagate with motion model sampling



PARTICLE FILTERS

○ Derivation (*why the measurement model is the weight assignment*):

- Consider the particles as state sequence samples

$$x_{(0:t)}^{[i]} = x_0^{[i]}, x_1^{[i]}, \dots, x_t^{[i]}$$

- Form belief over entire sequence

$$bel(x_{0:t}) = p(x_{0:t} | u_{1:t}, y_{1:t})$$

- Instead of just

$$bel(x_t) = p(x_t | u_{1:t}, y_{1:t})$$

- This is an enormous state to approximate with a set of particle, but no matter, for derivation only

PARTICLE FILTERS

○ Derivation

- Using Bayes Theorem, expand belief about last measurement

$$\begin{aligned} \text{bel}(x_{0:t}) &= p(x_{0:t} | u_{1:t}, y_{1:t}) \\ &= \eta p(y_t | x_{0:t}, u_{1:t}, y_{1:t-1}) p(x_{0:t} | u_{1:t}, y_{1:t-1}) \end{aligned}$$

- The Markov assumption remains valid

$$= \eta p(y_t | x_t) p(x_{0:t} | u_{1:t}, y_{1:t-1})$$



Normalizer

PARTICLE FILTERS

○ Derivation

- Conditional probability can be used to expand the last distribution

$$\begin{aligned} &bel(x_{0:t}) \\ &= \eta p(y_t|x_t) p(x_t|x_{0:(t-1)}, u_{1:t}, y_{1:t-1}) p(x_{0:t-1}|u_{1:t}, y_{1:t-1}) \end{aligned}$$

- Apply the Markov assumption again yields

$$bel(x_{0:t}) = \eta p(y_t|x_t) p(x_t|x_{t-1}, u_t) p(x_{0:t-1}|u_{1:t}, y_{1:t-1})$$

~ current belief at time t-1

Measurement model

Motion model

PARTICLE FILTERS

○ Derivation

- The sequence $x_{0:t-1}$ does not depend on u_t

$$\begin{aligned} bel(x_{0:t}) &= \eta p(y_t|x_t) p(x_t|x_{t-1}, u_t) p(x_{0:t-1}|u_{1:t-1}, y_{1:t-1}) \\ &= \eta p(y_t|x_t) p(x_t|x_{t-1}, u_t) bel(x_{0:t-1}) \end{aligned}$$

- Breaking into two steps
 - Prediction

$$\overline{bel}(x_{0:t}) = p(x_t|x_{t-1}, u_t) bel(x_{0:t-1})$$

- i^{th} particle generated by this distribution is an element of the predicted belief particle set

PARTICLE FILTERS

○ Derivation

- The measurement update uses importance sampling to generate a particle set representation of belief
 - Weighting, based on relation to predicted belief is

$$\begin{aligned}w_t^{[i]} &= \eta \frac{bel(x_{0:t})}{\overline{bel}(x_{0:t})} \\&= \frac{\eta p(y_t|x_t) p(x_t|x_{t-1}, u_t) bel(x_{0:t-1})}{p(x_t|x_{t-1}, u_t) bel(x_{0:t-1})} \\&= \eta p(y_t|x_t)\end{aligned}$$

- Which confirms the use of the measurement model as the weighting parameter
- This confirms that particles sets are distributed according to full belief sequences, which means must hold for current state too

PARTICLE FILTERS



- Example (*Single variable*)

- Returning to the temperature control problem
 - State is current temperature
 - One dimensional example
 - Prior: Uniform over temperature range
- **Motion Model:** Decaying temperature +furnace input + disturbances (opening doors, outside effects)

$$x_t = 0.8x_{t-1} + 3u_t + r_t$$

$$A = 0.8, \quad B = 3$$

$$r_t \sim N(0,2)$$

PARTICLE FILTERS



- Example

- Measurement Model:

- Directly measure the current temperature

$$y(t) = x(t) + \delta_t$$

$$\delta_t \sim N(0,4)$$

- Controller design:

- Bang bang control, based on current estimate of temperature

$$u(t) = \begin{cases} 1 & \mu_t < 2 \\ 0 & \mu_t > 10 \\ u(t-1) & \text{otherwise} \end{cases}$$

PARTICLE FILTERS



○ Particle filter calculations

1. Transform and Sample from Gaussian

$$\bar{x}_t^{[i]} \sim p\left(x_t \mid x_{t-1}^{[i]}, u_t\right) = 0.8x_{t-1}^{[i]} + 3u_t + r_t$$

2. Define weights from measurement model, with Gaussian noise centered at predicted measurement location

$$w_t^{[i]} = p(y_t \mid \bar{x}_t^{[i]})$$

3. Resample from predicted belief particle set using weights to generate belief particle set

PARTICLE FILTERS


- Resulting particle filter Matlab code



```
%Particle filter estimation
for i=1:I
    e = sqrt(R)*randn(1);
    Xp(i) = A*X(i) + B*u(t) + e;
    w(i) = normpdf(y(t),C*Xp(i),Q);
end
W = cumsum(w);
W = W/max(W);
for j=1:I
    i = find(W>rand(1),1);
    X(j) = Xp(i);
end
```


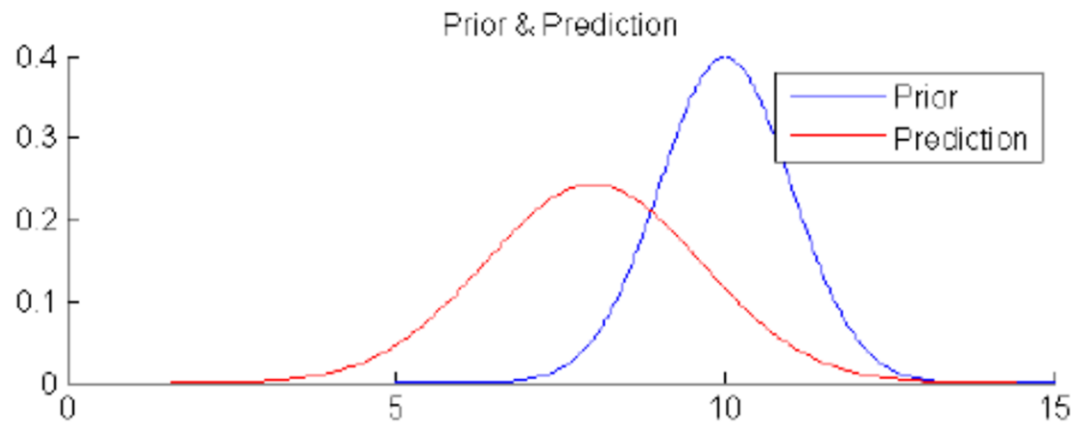
PARTICLE FILTERS

- Priors, comparing KF and particle filter:





PARTICLE FILTERS

○ Prediction:



PARTICLE FILTERS

○ Belief:

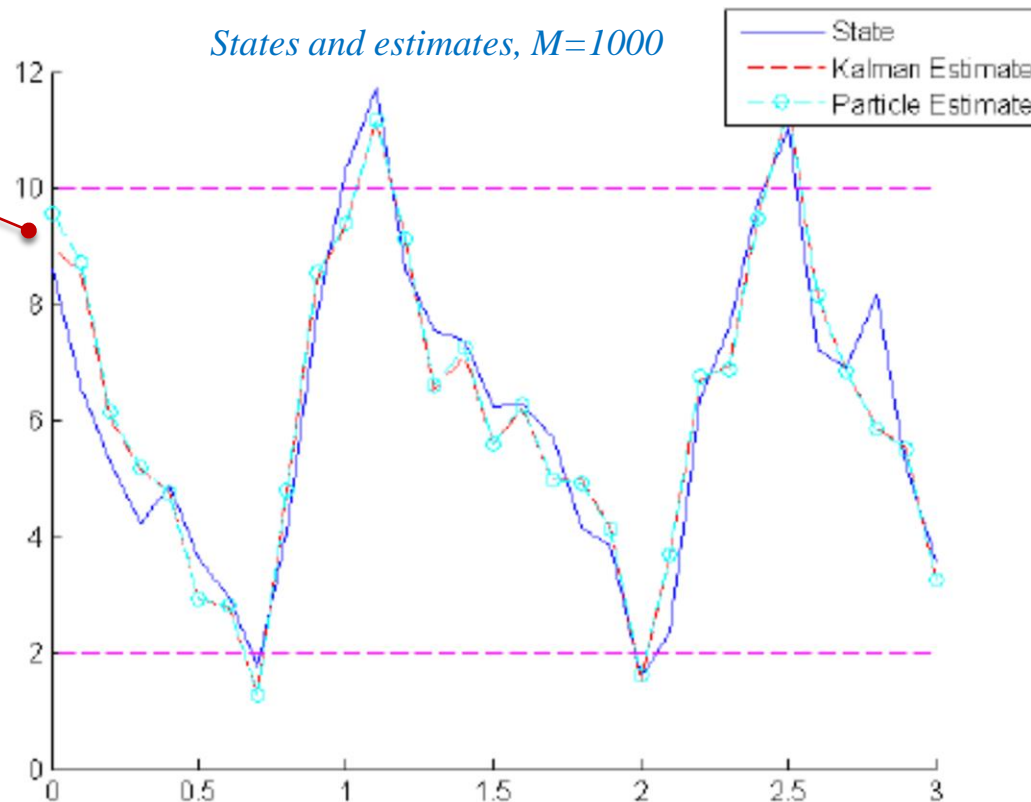


PARTICLE FILTERS

- Comparison of Gaussian parameters
 - KF vs PF (1000, 100, 10)

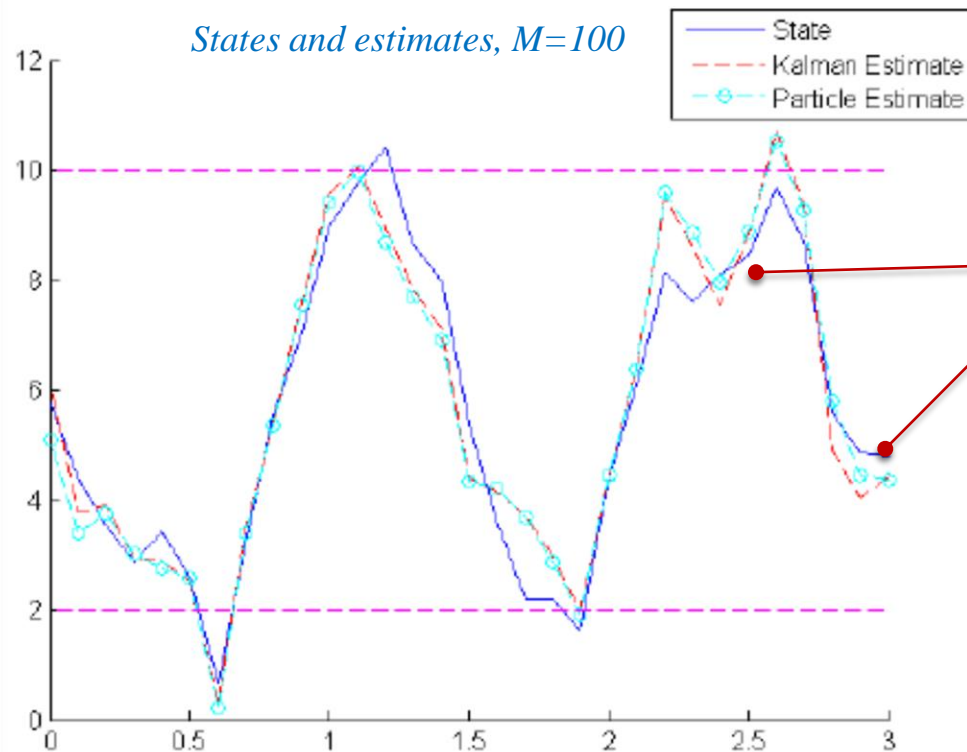


Some minor deviations based on the initial belief



PARTICLE FILTERS

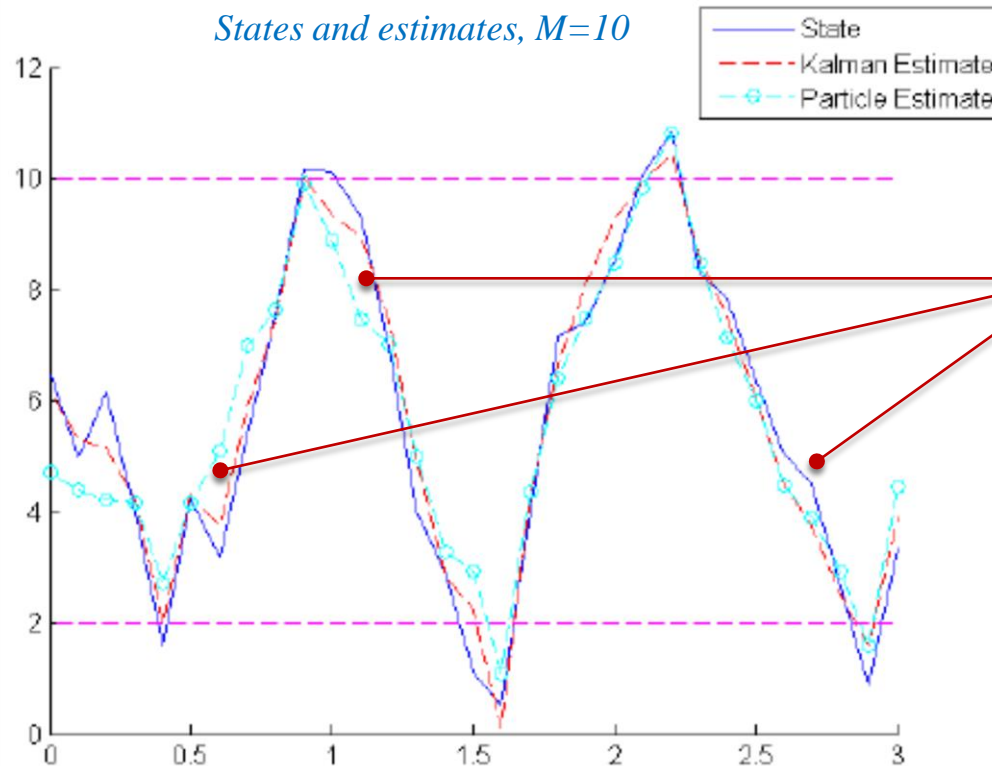
- Comparison of Gaussian parameters
 - KF vs PF (1000, 100, 10)



*More variation
relative to the KF*

PARTICLE FILTERS

- Comparison of Gaussian parameters
 - KF vs PF (1000, 100, 10)



*Significant
deviation w.r.t.
the KF*

PARTICLE FILTERS

- Comparison of run times (in Matlab):




Algorithm	Run (<i>computation</i>) Time
Kalman Filter	0.005164
Particle Filter – <i>10 particles</i>	0.043191
Particle Filter – <i>20 particles</i>	0.06965
Particle Filter – <i>100 particles</i>	0.2188
Particle Filter – <i>1000 particles</i>	1.8740

Particle filters: Due to high computational cost we only use them when we need them (in very complex problems).



PARTICLE FILTERS

- EKF/UKF/Particle showdown
 - Aircraft flyover example



PARTICLE FILTERS



○ Results (*Video*)



Aircraft flyover results using 500 particles (each representing a possible solution on where the aircraft is)

PARTICLE FILTERS

- Not always better, hard to tune



PARTICLE FILTERS

○ Interpreting particle sets

- In order to use a particle filter in robotics, we must somehow extract relevant information from particles that we can use to control robots.
- **Density Extraction**
 - Gaussian approximation:
 - Simply calculate mean and covariance of set
 - Only really useful for unimodal distributions
 - Used most often for control applications
 - K-means algorithm:
 - Approximate density with a mixture of K Gaussians
 - Requires clustering of particles
 - Kernel density estimation:
 - Use each particle as the center of a continuous kernel function
 - Add all kernels together to generate a pdf
 - Linear in the number of particles

PARTICLE FILTERS

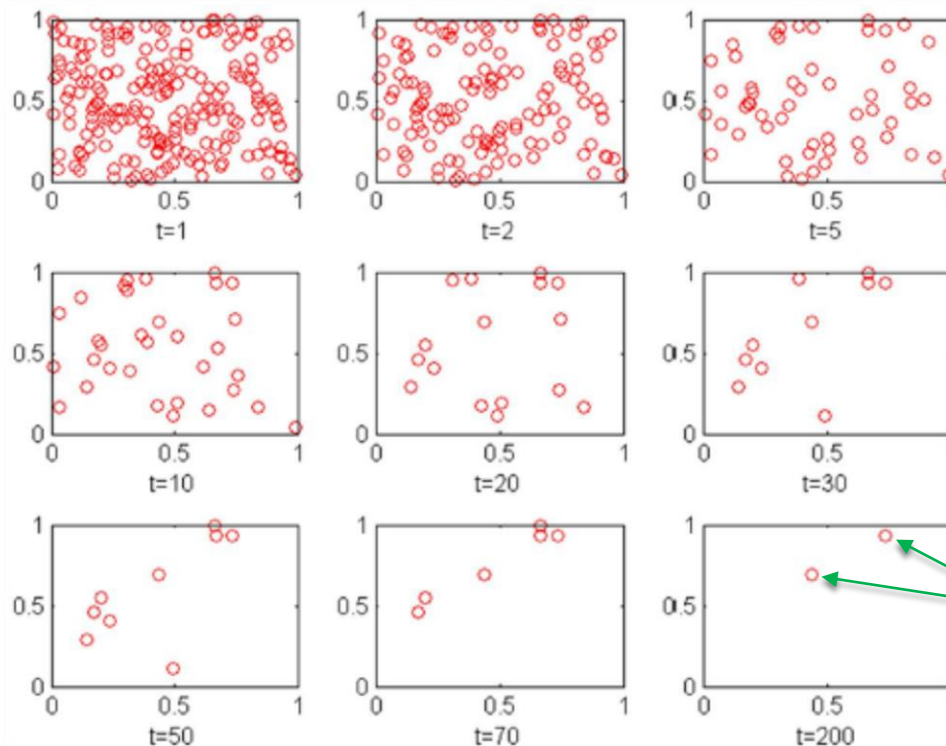
○ Sample Variance

- Since continuous distributions are approximated by a discrete set of samples, errors occur
- Each time a particle filter is run (with random sampling) a different particle set will result
- **Extreme case:**
 - No motion: $x_t = x_{t-1}$
 - No measurements, uniform weights on each particle
 - Uniform prior over 2D space
 - What will happen to the particle set as we update the particle filter?
 - Essentially repeating the resampling step with uniform weight on all particles

PARTICLE FILTERS

○ Example

- Particle deprivation (*the richness of the particle set is important to prevent deprivation*)



Each of these 2 particles include many copies, maintaining the size of the original sample set

PARTICLE FILTERS

- Excessive resampling can lead to particle deprivation
 - Motion sampling adds variety to particle set
 - Do not resample when no motion occurs
 - *Instead update weights multiplicatively for each measurement*
 - If problems arise
 - *Apply low variance sampling*
 - *Artificially disperse samples as well*
 - *Add random samples after resampling*
 - Referred to as variance reduction
 - *Reducing the variance in the particle set approximation*

PARTICLE FILTERS

○ **Summary:**

- Use particle sets instead of parameterizations to represent distributions
- Inherently an approximation, introduces errors
- Propagate samples through motion model by sampling from model distribution
- Weight samples using measurement probability given sample state as true state
- Define belief distribution through samples and weights (particles) or post resampling
- Many extensions, nuances, issues, advanced techniques